

Patrick Ziegler

Design and Evaluation of a Network Softwarization Platform for Edge and In-Network Computing

Diploma Thesis Defense, 24.10.2019

Deutsche Telekom Chair for Communication Networks


Motivation



- Motivation
 - **Dynamically deploy functions at the network edge**
 - Enable new use cases for IoT, Mobile Edge Cloud etc.
 - Offloading computation tasks (also for cooperation between clients)
- This is supported by **Network Function Virtualization (NFV)**
- State of the Art NFV platforms are
 - Based on **complex** cloud infrastructure platforms like OpenStack
 - Mainly focused on **business related problems**
 - Restricted to **fixed deployments**

Problem statement


- Cloud platforms provide **abstracted view** of infrastructure
 - This **hides network related information** from NFV platforms
- Edge deployments should provide minimal latency
 - NFV needs **latency and topology awareness**
- Manual optimization is too complex
 - NFV should provide **autonomous optimization and deployment**



Problem statement

- Cloud platforms provide **abstracted view** of infrastructure
 - This **hides network related information** from NFV platforms
- Edge deployments should provide minimal latency
 - NFV needs **latency and topology awareness**
- Manual optimization is too complex
 - NFV should provide **autonomous optimization and deployment**

Goals of this work



Scope of this work


- Design and implement a lightweight NFV platform from scratch
 - Serves as **proof of concept** for new ideas
 - May arise **new insights** on NFV platforms („Wissen kommt von machen“)
 - **„Instant NFV on your laptop“** → good for research
- Three problem fields identified
 - **Monitoring** (Latency, Topology, Resources)
 - **Management and Orchestration**
 - **Placement** (Algorithms)
- These combined constitute a **new adaptive approach to NFV**



Monitoring




Problem setting




- How to measure link latency in OpenFlow networks?
- OpenFlow only supports **predefined set of operations**
 - Not (easily) possible to implement new functionality
- State of the art monitoring solutions use **timestamped probing packets** injected by SDN controller
 - Compensation of **delay introduced in control network** necessary
 - How to handle links between switches and **compute nodes**?

Background: eBPF / XDP


- Extended Berkeley Packet Filter (eBPF)
- A **virtual machine** inside the Linux kernel
 - Resembling the Java Virtual Machine
- eBPF Programs
 - are **attached to provided hooks** and executed when events occur
 - have to **pass a verifier** before being allowed to be attached to a hook
 - can **manipulate in- and outgoing packets**
 - decide over packets further processing (**PASS, DROP, REDIRECT**)




1. Controller sends probing packets to registered switches



2. Switches attempt to do LLDP Looping with all adjacent nodes



3. Adjacent nodes send probing results back to Controller




Review: LLDP Monitoring

- XDP allows to return packet **without incurring queueing delays**
 - this enables **seperation of individual delays**
- **No clock synchronization** needed
- Transmission delay is assumed to be symmetric

Negligible

$$\text{Latency} = \cancel{T_{Transmission}} + T_{Propagation} + T_{Queueing} = RTTQ - RTT/2$$

Symmetric




Precision on empty link

Reference	50 ms
Mean value	50.030141 ms
Std. dev.	0.009187 ms


Management and Orchestration



Problem setting




Overview of involved components




Background: The Serverless Paradigm

- Core idea
 - Transfer the **responsibility for** infrastructure **provision and deployment** from the administrator **to the platform**
- Deployment units are *functions*
 - typically **stateless**
 - running inside containers or virtual machines
 - functions are **handling full HTTP requests** in conventional platforms
- Provision **as the need arises**
 - based on incoming requests
 - with **automatic scaling** under consideration of resource usage (CPU)
- Concept most effective when **request rate shows high volatility**


1. Peer requests a network service




2. NFVO implants service hook




3. Traffic hits implanted hook




4. NFVO requests VNF deployment




5. Local forwarding rules implanted




6. Global forwarding rules implanted




7. Packet forwarded to first VNF




8. Packet visits succeeding VNFs




9. Packet reaches destination




10. Network service fully established



11. Flow removal due to idle timeout




12. VNFs are automatically removed



Review: Serverless MANO

- The developed **serverless network service lifecycle**
 - Is based on OpenFlow and label based switching
 - Reduces the amount of **manual administration to zero**
 - Keeps **resource usage as low as possible**
- Startup delay of network services is typically > 100 ms
 - Drawback that has to be considered when designing services




Short demo at the end of the presentation

Placement




Hasty Traveller Algorithm: Concept

- Design Goal
 - find shortest path **visiting a defined number of workers** along the way (chain length)
 - Travelling Salesman problem
- **Recursive search** in mesh representation
 - at each node, visit adjacent nodes and sum up incurred latency
 - to **limit the complexity**, only visit some adjacent nodes with lowest link latency
- From all visited paths, **the one with lowest latency is chosen**





Hasty Traveller Algorithm: Expected Latency

- Expected Latency Plot
 - Hasty Traveller with various limit values (**HT-limit**)
 - compared to 1000 random placements (**RP-1000**)
- Lower boundary of grey area represents best placement
- Upper boundary of grey area represents worst placement
- Results naturally depend on the underlay topology



Hasty Traveller Algorithm: Timing

- Plots show time it took to reach a solution
- With increasing limit, Hasty Traveller becomes very expensive
- At some point, choosing the best of 1000 random placements is faster
- A limit of 3 shows **reasonable complexity and good results** for typical chain lengths



Conclusion



Contributions

- State of the art NFV platforms
 - are missing awareness of the underlay network
 - can therefore not provide latency optimized deployments
- This work introduced new concepts to address these issues
 - **eBPF based integrated LLDP Monitoring** → increases network awareness
 - **Serverless network service lifecycle** → reduces administration complexity
 - **Graph theoretic placement strategy** → supports latency optimized deployments
- These ideas enabled the **design of a lightweight NFV platform**
 - Proof of concept implementation can be accessed via
<https://git.comnets.net/patrickziegler/nfv-mano>

Outlook

- The developed system lays the **groundwork for further research** on
 - **Resource allocation** and **VNF benchmarking** for predicting computing delays
 - Early adoption of recently proposed **Network Service Headers** (NSH)
 - Evaluating the operation of P4 programmable switches against the OpenFlow based approach used in this work
- The proof of concept implementation
 - can be further developed to serve as a platform for **teaching NFV**
 - serves as a **reproducible environment** for research on new network services


Live Demo


Patrick Ziegler


Diploma Thesis Defense, 24.10.2019

Deutsche Telekom Chair for Communication Networks

Monitoring system supports additional probing on Workers










Image source:

S. K.Mohapatra, J. Bhuyan, and H. N.Narang, "Planning and managing virtualized next generation networks," International journal of Computer Networks & Communications, vol. 7, pp. 01–16, 11 2015

